Neuro-Symbolic and Cognitive AI Chatbot for Adaptive Stress Regulation Among Undergraduates

25-26J-138

Proposal Report Draft

F N F Ramla

BSc (Hons) Information Technology Specializing in Data Science

Department of Computer Science

Sri Lanka Institute of Information Technology

Sri Lanka

August 2025

Neuro-Symbolic and Cognitive AI Chatbot for Adaptive Stress Regulation Among Undergraduates

Neuro Symbolic Emotion Reasoning Engine

25-26J-138

Proposal Report Draft

F N F Ramla

BSc (Hons) Information Technology Specializing in Data Science

Department of Computer Science

Sri Lanka Institute of Information Technology

Sri Lanka

August 2025

Declaration

We declare that this is our own work, and this proposal does not incorporate without acknowledgement any material previously submitted for a degree or diploma in any other university or Institute of higher learning and to the best of our knowledge and belief it does not contain any material previously published or written by another person except where the acknowledgement is made in the text.

Signature		**Date**	**15/08/2025**
Signature of Supervisor		**Date**	**28/08/2025**
Signature of Co-Supervisor		**Date**	**28/08/2025**

Abstract

This proposal presents a research project focused on designing, developing, and evaluating a Neuro-Symbolic Emotion Reasoning Engine (NSERE) to enhance adaptive stress regulation for university students. The project addresses a critical gap in digital mental health interventions, where existing solutions: from rule-based chatbots to general-purpose Large Language Models (LLMs) lack the transparency, personalization, and psychological depth required for effective, high-stakes applications. The core problem is the "black-box" nature of neural networks and the rigidity of symbolic systems, which lead to generic or unsafe responses. To solve this, the research proposes a hybrid neuro-symbolic architecture that combines the strengths of deep learning for emotional pattern recognition with the explainable, logical reasoning of psychological knowledge graphs. The methodology outlines a phased approach, including data acquisition, hybrid model development, and a rigorous validation plan. The anticipated outcome is a validated prototype that provides transparent, psychologically grounded emotional support, with significant potential for commercialization by licensing the technology to universities.

Contents

Declaration 3
Abstract 4
List of Figures 6
List of Tables 6
List of Abbreviations 6
1 Introduction 7
2 Background and Literature Survey 9
3 Research Gap 12
4 Research Problem 14
5 Objectives 15
5.1 Main Objectives 15
5.2 Specific Objectives 15
6 Methodology 16
6.1 Tasks and Sub-Tasks 16
6.2 System Architecture Diagram 18
6.2.1 Overall System Architecture Diagram 18
6.2.2 Component System Architecture Diagram 18
6.3 Work Breakdown Structure 19
6.4 Gantt Chart 20
6.5 Required Materials 20
6.6 Anticipated Conclusion and Real-World Application 21
7 Project Requirements 22
7.1 Functional Requirements 22
7.2 Non-Functional Requirements 23
7.3 User Requirements 24
7.4 Use Cases (Tentative) 24
7.5 Test Cases (Tentative) 25
8 Commercialization and Entrepreneurship Potential 26
9 Budget and Budget Justifications 29
10 References 30
11 Appendices 32

List of Figures

Figure 6.1 Overall System Architecture Diagram 18
Figure 6.2 System Architecture of the Neuro-Symbolic Emotion Reasoning Engine 18
Figure 6.3 Work Breakdown Strcuture 19
Figure 6.4 Gantt Chart 20

List of Tables

Table 8.1 Competitive Landscape & Feature Comparison 28
Table 9.1 Budget and Budget Justification 29

List of Abbreviations

Abbreviation	**Definition**
NSERE	Neuro Symbolic Emotion Reasoning Engine
UI	User Interface
AI	Artificial Intelligence
NLP	Natural Language Processing
LLM	Large Language Model

1 Introduction

University students face escalating mental health challenges, including heightened stress, anxiety, and depression, which significantly impair academic performance and overall well-being[1]. The inherent pressures of higher education, such as rigorous curriculum, demanding evaluations, and peer comparisons, contribute substantially to these psychosocial risks[1]. The pervasive nature of these issues necessitates scalable and accessible mental health support. Digital mental health interventions, including mobile applications and chatbots, offer a promising avenue for stress management among young adults by enhancing accessibility and potentially mitigating healthcare disparities[2]. The current student-to-counselor ratios often fall short of recommended standards, leading to prolonged waiting times for in-person consultations and suboptimal care, underscoring the critical need for advanced digital solutions to bridge this systemic resource gap[3].

Existing mental health chatbots predominantly rely on Natural Language Processing (NLP) for text input and rule-based programming to deliver prewritten content, with limited integration of advanced generative AI models.[4] While these systems offer accessibility and interactivity, they frequently suffer from repetitive content, a limited understanding of human language, and insufficient personalization.[4] General-purpose Large Language Models (LLMs), despite demonstrating capabilities in cognitive reframing and affect recognition, often provide generic advice, lack contextual inquiry, and are currently deemed unsuitable as primary therapeutic agents, particularly in crisis scenarios where they fail to perform adequate risk assessments.[5]Furthermore, these models are often perceived as lacking profound emotional understanding and empathy, which are crucial for effective therapeutic engagement.

The core limitation stems from the architectural dichotomy of current AI paradigms. Purely neural models, while adept at pattern recognition, function as "black-boxes," hindering transparency and interpretability in sensitive mental health applications.[6]Conversely, traditional symbolic AI, despite offering inherent interpretability and precise reasoning, struggles with processing unstructured data and maintaining scalability in complex, dynamic environments.[7]This fundamental architectural constraint prevents current AI from achieving the nuanced, trustworthy emotional reasoning required for comprehensive stress regulation.

To address these critical limitations, this proposal introduces a Neuro-Symbolic Emotion Reasoning Engine. This engine is designed to integrate symbolic AI with deep learning to emulate human-like reasoning, enabling precise identification and appropriate response to stress-inducing conversational triggers. Neuro-symbolic AI represents a transformative approach, combining the strengths of statistical AI (machine learning) with human-like symbolic knowledge and reasoning to pursue artificial general intelligence. This hybrid methodology is uniquely positioned to overcome the inherent weaknesses of purely neural or symbolic systems, enhancing generalization, reasoning capabilities, and scalability while simultaneously improving transparency and data efficiency.[8] By combining raw conversational data with formalized psychological and clinical knowledge, the engine aims to support transparent and explainable diagnostic and intervention systems in mental health contexts.[9] This approach signifies a

qualitative leap from mere pattern-matching to a genuine cognitive partnership, offering a targeted solution to fundamental AI challenges in this sensitive domain.

The proposed Neuro-Symbolic Emotion Reasoning Engine leverages neuro-symbolic hybrid models to extract mental state inferences by integrating language embeddings with psychological knowledge graphs. Neuro-symbolic AI fundamentally combines neural networks, which excel at perception and pattern recognition, with symbolic AI, which provides structured knowledge, logic, and rule-based reasoning.[8]

Emotion recognition from text is achieved through advanced NLP techniques, including deep learning models, which analyze linguistic indicators such as first-person pronouns and emotional terms to infer mental states.[10] Emotion embedding models further leverage convolutional neural networks to extract nuanced emotional content from text.[25] The critical novelty lies in the synergistic integration of these language embeddings with psychological knowledge graphs. Knowledge graphs provide structured representations of entities and their interrelations, holding significant potential in psychiatric research and mental health exploration by incorporating multi-dimensional data and adapting to emerging knowledge.[11]This combination allows the neural component to perceive subtle linguistic cues, while the symbolic knowledge graphs provides the contextual and psychological rules necessary to interpret these cues into meaningful mental states, enabling sophisticated.

The emphasis on explainability is central to this approach. Neuro-symbolic AI inherently aims for transparent outcomes by leveraging the clarity of symbolic learning.[12] This intrinsic transparency, rather than post-hoc explanations, is vital for building trust and accountability in sensitive mental health applications.[9] Furthermore, the integration of psychological knowledge graphs, particularly those incorporating cultural and contextual data, directly addresses the limitations of algorithmic bias and generic responses prevalent in purely neural models. This ensures that the chatbot provides personalized and culturally sensitive emotional support, moving beyond simple pattern matching to deliver psychologically grounded and ethically robust interventions.

2 Background and Literature Survey

The mental health of university students is a growing public health concern, with a significant prevalence of stress, anxiety, and depression impacting academic performance and overall well-being.[1] Modern students often exhibit reduced stress tolerance and inadequate coping mechanisms compared to previous generations, exacerbating these issues. Traditional mental health support systems in universities face severe limitations, primarily due to unfavorable student-to-counselor ratios, leading to extended waiting periods, suboptimal care[3] and the stigma surrounding mental health wellness adds up to this. This systemic gap underscores the urgent need for scalable and accessible interventions.

In response, digital mental health interventions including mobile applications and chatbots, have emerged as promising tools for stress management among young adults.[2] These AI-driven tools offer immediate, accessible, and often anonymized support, which can reduce the stigma associated with traditional mental healthcare and overcome barriers to professional care [13].Digital mental health interventions have shown potential in alleviating symptoms of stress, anxiety, and depression, their effectiveness and sustained usage depend on factors such as user-friendliness, personalization, and cultural considerations [2] .Conversational AI, in the form of therapeutic chatbots, has emerged as a promising avenue for providing immediate, accessible, and consistent support.[5] However, the current generation of these systems, predominantly built on traditional machine learning (ML) or large language models (LLMs), has shown significant limitations in handling the complexity and nuance required for genuine psychological support.[5] These models often fail to capture the context and causality behind emotional states, relying instead on surface-level pattern matching that can lead to misinterpretations or inappropriate responses.

Neuro-symbolic AI represents a transformative hybrid approach that addresses the limitations of both purely symbolic and purely neural AI systems by combining their respective strengths.[8]Traditional neural networks (NNs) excel at learning patterns from massive, unstructured datasets without requiring pre-coded rules.[8] However, this strength comes with several well-documented disadvantages: they require vast amounts of labeled data, struggle with logical reasoning and generalization beyond their training data, and operate as opaque "black boxes" that cannot explain their decisions,[8] Conversely, symbolic AI systems are adept at handling explicit rules, structured knowledge, and logical reasoning, providing inherent transparency.[8] Their rigidity, however, makes them ill-suited for learning from unstructured data or adapting to new information without manual intervention.

By merging these paradigms, neuro-symbolic AI aspires to embody two fundamental aspects of human cognition: the ability to learn from experience (neural) and the capacity to reason with acquired knowledge (symbolic). [8]This concept is often likened to Daniel Kahneman's dual-process theory of reasoning, which distinguishes between a fast, intuitive System 1 and a slower, deliberate, and logical System 2.[8] In the context of the proposed engine, the neural component (System 1) would handle the rapid, unconscious pattern recognition of language and emotional cues, while the symbolic component (System 2) would provide the slow, explicit reasoning required for tasks like personalized intervention planning.This framework provides a profound theoretical underpinning, elevating the engine from a simple technical solution to a cognitively plausible model.

While some researchers debate whether the increasing logical capabilities of LLMs could obviate the need for a separate symbolic layer , other sources highlight a fundamental inadequacy of LLMs as standalone conversational agents for high-stakes applications. The debate is not about whether LLMs can perform some reasoning, but whether they can do so reliably and transparently. Current literature on therapeutic chatbots and high-stakes financial applications points to the "black box problem" and the risk of "hallucinations" the generation of factually incorrect or inappropriate content as core limitations of purely neural models.[4] This suggests that even as LLMs improve, the need for a separate, auditable symbolic layer for safety, explainability, and policy adherence remains paramount, especially in a sensitive domain like mental health. The proposed engine's neuro-symbolic design directly addresses this ongoing industry "blind spot".

The landscape of natural language processing (NLP) for emotional inference has evolved significantly. Early approaches relied on rule-based systems using predefined lexicons and grammatical rules to identify emotions.[10] With the advent of deep learning, the field shifted to sophisticated models such as Convolutional Neural Networks (CNNs) and Recurrent Neural Networks (RNNs) that learn to classify emotions from text embeddings.[10] These models have achieved high accuracy on simple emotion detection tasks, which involve analyzing text to determine its emotional tone or sentiment.[10]

However, this review identifies a critical semantic and cognitive discrepancy between **emotion detection** and **emotion reasoning**. Emotion detection is a classification task that identifies a sentiment or emotion (e.g., "sadness" or "anger") from text. Emotion reasoning, by contrast, is a higher-order cognitive function that involves understanding the *causes, context, and implications* of that emotion. For example, a deep learning model might correlate the phrase "my boss gave me more work" with "stress" based on statistical patterns. However, it lacks the explicit, causal knowledge to understand

why this phrase is a stressor.

The literature indicates that this is a foundational gap in current approaches. Existing knowledge graphs, which form the basis for symbolic knowledge, have focused on general factual knowledge and have "significant deficiencies in the understanding and reasoning of emotional knowledge". The path toward true emotion reasoning requires a shift from simple correlation to a structured understanding of causality.

To bridge this gap, symbolic knowledge graphs have been combined with neural embeddings.

The application of AI in mental health is a high-stakes domain where reliability, transparency, and ethical safeguards are non-negotiable. A comparative analysis of neuro-symbolic systems versus purely neural or statistical models reveals significant advantages for the hybrid approach in this context.

Neuro-Symbolic vs. Purely Neural/Statistical Models

- **Effectiveness and Accuracy:** Neuro-symbolic systems have demonstrated superior performance in tasks that require logical constraints and robustness to data errors. In mental health, this translates to enhanced diagnostic accuracy and the ability to generate more reliable, personalized therapeutic recommendations.[9]
- **Explainability and Trust:** One of the most critical advantages of neuro-symbolic AI is its enhanced explainability.[8] Unlike opaque neural models, symbolic systems can provide human-interpretable "reasoning chains" and explicit logic for their decisions. This transparency is essential for building clinician trust and for ethical deployment, as it allows mental health professionals to validate AI-generated insights against established psychiatric frameworks and ensure clinically sound decision-making.[14]
- **Data Efficiency and Robustness:** Purely neural models have a high requirement for data, making them ineffective in data-scarce domains like rare mental health conditions.[9] neuro-symbolic AI systems mitigate this limitation by leveraging pre-coded rules and structured knowledge, allowing them to generalize beyond limited training data and operate effectively in data-costly environments.[8]

The performance of existing therapeutic chatbots (e.g., Wysa, WoeBot) and general-purpose LLMs (e.g., GPT-4) provides crucial context. Research indicates that while general-purpose LLMs can sometimes outperform therapeutic bots in rectifying cognitive biases and adapting to emotional nuances, they still have significant limitations.[5] The performance of general-purpose LLMs, which are not designed for therapeutic contexts, can be attributed to their sheer scale and fluency, which allows for a broader flexibility in handling complex linguistic patterns.[5] However, this finding does not negate the need for a specialized neuro-symbolic engine; rather, it highlights a different kind of gap. Current therapeutic bots may be "brittle," relying on rigid, pre-scripted conversational flows.[4] Their primary limitation is not a lack of fine-tuning, but an absence of a flexible, adaptive reasoning core. The general-purpose LLMs demonstrate that adaptability is key. This implies that the optimal solution is not a more fine-tuned, static therapeutic bot, but a system that can combine the adaptability and fluency of neural networks with a structured, auditable reasoning layer.This combination is the very definition of neuro-symbolic AI.

Furthermore, the literature highlights critical deficiencies in all current chatbots, including their lack of genuine empathy and intuition, risk of oversimplifying complex issues, potential for harmful and unsafe responses (hallucinations), and inability to provide a predictable audit trail or adhere to clinical policies.[4] These deficiencies underscore the need for a fundamental architectural shift. This framework is not merely a feature but a fundamental design philosophy for mental health AI. The goal is to create a transparent, reliable tool that therapists can trust and use to augment their own judgment, making the AI a powerful assistant rather than a "digital therapist" that risks undermining professional judgment or causing harm.

3 Research Gap

Despite significant advancements, several critical gaps persist in the literature that justify the need for the proposed Neuro-Symbolic Emotion Reasoning Engine

1. **Lack of Explainability in Deep Learning Models:** While LLMs excel at pattern recognition and generating coherent responses, their "black-box" nature fundamentally limits transparency and interpretability [15]In sensitive domains like mental health, where trust, accountability, and clinical validation are paramount, a system that cannot articulate
   why it inferred a particular emotional state or suggested an intervention is severely constrained Existing neuro-symbolic AI research on explainability acknowledges that simply combining neural and symbolic components does not automatically guarantee full transparency, and further work is needed to bridge the gap between technical explanations and human understanding [15]

2. **Insufficient Psychological Grounding and Nuanced Reasoning:** Purely data-driven models, despite their performance in emotion detection, often lack genuine self-awareness, personal experience, and the capacity for subjective feeling, which are crucial for deep emotional understanding and authentic empathy [16]They struggle with complex, multi-step reasoning and commonsense understanding, leading to generic or even absurd responses in ambiguous situations.While psychological knowledge graphs exist, their deep integration with real-time conversational language embeddings for nuanced, context-aware emotional inference remains an area for significant advancement.

3. **Limitations in Real-time Adaptive Stress Regulation:** Current mental health chatbots, whether rule-based or LLM-driven, often provide prewritten or generic content, lacking the dynamic adaptability required for real-time, personalized stress regulation [4]The ability to accurately identify stress-inducing triggers and infer mental states with human-like cognitive processes in real-time conversational flow is still a challenge.

This research directly addresses these gaps by proposing the Neuro-Symbolic Emotion Reasoning Engine. This approach will advance the state of the art by:

- **Enabling Explainable Emotional Intelligence:** By leveraging neuro-symbolic hybrid models,the engine will integrate structured symbolic reasoning with probabilistic inference to provide transparent, human-interpretable explanations for its emotional inferences and suggested interventions [12]This moves beyond opaque "black-box" models, fostering crucial trust with users and clinicians.

- **Achieving Deep Psychological Grounding:** The engine will extract mental state inferences by combining language embeddings with comprehensive psychological knowledge graphs.This fusion will allow the chatbot to move beyond superficial pattern matching to a semantically rich, context-aware, and psychologically informed understanding of emotional states and their triggers, enabling reasoning ,

- **Facilitating Human-like Reasoning for Adaptive Stress Regulation:** By integrating the strengths of neural perception and symbolic reasoning, the engine will simulate human-like cognitive processes for identifying and responding to stress-inducing cues [17]This will enable the chatbot to provide more accurate, empathetic, and therapeutically relevant responses, laying a foundational layer for adaptive stress regulation that is currently lacking in existing solutions.

4 Research Problem

Despite the growing need for accessible mental health support among university students, existing AI-driven chatbots face significant limitations. Current systems, whether rule-based or purely deep learning models, often lack genuine emotional understanding, empathy, and the ability to provide explainable reasoning. Purely neural models operate as "black-boxes," hindering user trust and clinical validation, while symbolic systems struggle with the dynamic, unstructured nature of human conversation and scalability.This results in generic, repetitive, or even unsafe responses, particularly in crisis situations, and a failure to offer truly personalized or psychologically grounded interventions.

These limitations have severe negative consequences: students may disengage from interventions due to a perceived lack of understanding or empathy, and the absence of transparency prevents clinicians from validating AI-driven recommendations, raising critical safety and ethical concerns. Given the escalating mental health crisis among university students and the shortage of human counselors, there is an urgent and timely need for AI solutions that can provide effective, trustworthy, and adaptive support. Solving this problem requires moving beyond current paradigms to a neuro-symbolic approach that integrates robust emotional perception with transparent, human-like reasoning, ensuring both efficacy and accountability in digital mental health interventions.

5 Objectives

5.1 Main Objectives

The main objective of this research is to **design, develop, and evaluate** an integrated **Neuro-Symbolic and Cognitive AI Chatbot** that delivers **adaptive stress regulation** for university students by combining four interdependent components: (1) a **Neuro-Symbolic Emotion Reasoning Engine** that fuses neural language understanding with symbolic psychological knowledge for explainable emotion inference, (2) **Cognitive Load Detection** from passive behavioral signals to assess mental strain in real time, (3) **Continual Learning** to adapt to evolving emotional patterns without catastrophic forgetting, and (4) an **Adaptive Multimodal Intervention Selector** to determine optimal intervention timing and modality. This objective is **Specific** in targeting explainable, context-aware emotional support; **Measurable** through predefined metrics including accuracy in emotion recognition, precision in cognitive load detection, and positive user satisfaction in pilot trials; **Achievable** within available datasets, computational resources, and expertise; **Relevant** to the urgent need for scalable and psychologically grounded mental health interventions for students; and **Time-bound** to be completed, tested, and validated within a **10-month research period**.

5.2 Specific Objectives

- **Neuro-Symbolic Architecture Design & Development** To design and develop a hybrid neuro-symbolic architecture that accurately identifies and classifies core emotional states from user conversational data with a classification good accuracy of bridging the gap between neural network-based pattern recognition and symbolic-based logical reasoning.
- **Knowledge Graph Construction & Reasoning Framework Development** To construct a formal knowledge graph representing psychological theories and emotional relationships to provide an explicit, explainable reasoning framework, with the goal of correctly interpreting the root causes of identified emotional states in over 90% of test cases.
- **Real-time Integration with Conversational Interface** To integrate the Neuro-Symbolic Emotion Reasoning Engine with the conversational interface to process real-time user input and generate emotionally-aware responses, with low latency.
- **System Validation & Explainability Evaluation** To establish a robust validation protocol using a curated dataset of university student stress dialogues, and to test the engine's explainability by generating human-readable reasoning traces that are rated well for clarity and coherence by independent evaluators.
- **Iterative Refinement and Performance Optimization** To iteratively refine the symbolic rule set and neural network model based on feedback from a pilot study, achieving a reduction in false-positive emotion classifications and an increase in user-reported empathy scores compared to a baseline model.

6 Methodology

The project will develop a reasoning engine that integrates **neural networks** for language understanding with **symbolic AI** for logical reasoning. This hybrid system will allow the chatbot to not only identify stress-inducing triggers in student conversations but also explain why a certain response is suggested. The process will involve dataset preparation, knowledge representation, hybrid model development, training and validation, and deployment into a chatbot prototype.

The project will employ a **research design** that integrates quantitative and qualitative approaches. The quantitative arm will focus on developing and validating a deep learning model for emotion classification from text, while the qualitative arm will capture the nuanced, context-dependent nature of student stress to inform a symbolic reasoning engine. The integration of these two methods is central to the neuro-symbolic approach.

The project will follow clear phases that build upon one another:

1. Phase 1: Project Initiation & Foundational Research: Establish the project's fundamental purpose and constraints
2. Phase 2: Data Acquisition and Preprocessing: Focus on collecting and preparing the necessary data and computational resources.
3. Phase 3: Core Development (Neuro-Symbolic System): Simultaneously develop the neural and symbolic components before integrating them.
4. Phase 4: Integration & Prototyping: Combine the developed components into a functional prototype.
5. Phase 5: Evaluation & Refinement: Rigorously test the system's performance and gather user feedback for improvements.

This model is chosen for its structured nature, ensuring that each phase is completed and validated before moving to the next, which is crucial for the complexity of a neuro-symbolic system.

6.1 Tasks and Sub-Tasks

Task 1: Project Initiation and Planning

- Sub-task 1.1: Research Problem Identification
  - Action: Define the core research problem (adaptive stress regulation in university students).
  - Output: Clear problem statement and justification.
- Sub-task 1.2: Literature Review & Gap Analysis
  - Action: Review existing work on chatbots, neuro-symbolic reasoning, and mental health AI.
  - Output: Identification of research gaps and novelty.
- Sub-task 1.3: Objectives & Scope Definition
  - Action: Define research objectives, expected contributions, and boundaries.
  - Output: Refined scope and measurable objectives.
- Sub-task 1.4: Project Planning & Milestones
  - Action: Develop timeline, allocate tasks, and define success criteria.

    - Output: Gantt chart / project plan.

Task 2: Data Acquisition and Preprocessing

- Sub-task 2.1: Public Dataset Acquisition
    - Action: Secure and download annotated datasets for emotion recognition.
    - Specific Datasets:
        - Mental Health Counseling Conversations (Kaggle).
        - ConvCounsel (student-focused, with emotion & counselor labels).
- Sub-task 2.2: Pilot Study for Student-Specific Data
    - Action: Conduct semi-structured interviews with 20 students (20 min each).
    - Output: Conversational transcripts focusing on stressful academic events.
- Sub-task 2.3: Data Preprocessing and Annotation
    - Action: Clean, normalize, and annotate data.
    - Processes: Remove PII, correct spelling, tokenize, lowercase, annotate emotions (e.g., Anxiety, Overwhelmed, Frustration, Sadness).

Task 3: Core Development – Neuro-Symbolic Model Development

- Sub-task 3.1: Neural Model Training for Emotion Classification
    - Action: Fine-tune a Transformer-based model for multi-class emotion classification.
    - Implementation: PyTorch and Hugging Face Transformers; evaluate with precision, recall, F1-score.
- Sub-task 3.2: Symbolic Knowledge Base Construction
    - Action: Build a Knowledge Graph with psychological concepts.
    - Structure: Nodes = Emotions, Triggers, Coping Strategies; Edges = relationships.
    - Rules: Define Production Rules using a rule engine.
- Sub-task 3.3: Hybrid Architecture Integration
    - Action: Pipeline where neural model outputs feed symbolic reasoning engine.
    - Process: Neural predictions → Confidence-weighted symbolic reasoning → Response strategy selection.

Task 4: System Integration and Interfacing

- Sub-task 4.1: Module Integration
    - Action: Integrate Neuro-Symbolic Emotion Reasoning Engine with other chatbot components:
        - Cognitive Load Detection Module
        - Continual Learning Module
        - Adaptive Stress Intervention Selector
- Sub-task 4.2: API Development
    - Action: Build RESTful APIs (Flask/Django).
    - Function: Expose NSERE services to the overall chatbot framework.
- Sub-task 4.3: Inter-Component Data Flow Testing
    - Action: Ensure seamless exchange of embeddings, reasoning outputs, and feedback signals between modules.
- Sub-task 4.4: Prototype Implementation
    - Action: Develop a basic front-end chatbot interface connected to all integrated

components.
    - Goal: Demonstrate end-to-end pipeline from input → Neuro-Symbolic Emotion Reasoning Engine → integrated reasoning → response.

Task 5: Evaluation and Refinement

- Sub-task 5.1: Technical Evaluation
    - Metrics: Accuracy, precision, recall, F1-score for neural model.
    - Symbolic Evaluation: Expert review & logic testing.
- Sub-task 5.2: Iterative Refinement
    - Action: Improve both neural and symbolic components.
    - Output: Optimized performance and user satisfaction.

6.2 System Architecture Diagram

6.2.1 Overall System Architecture Diagram

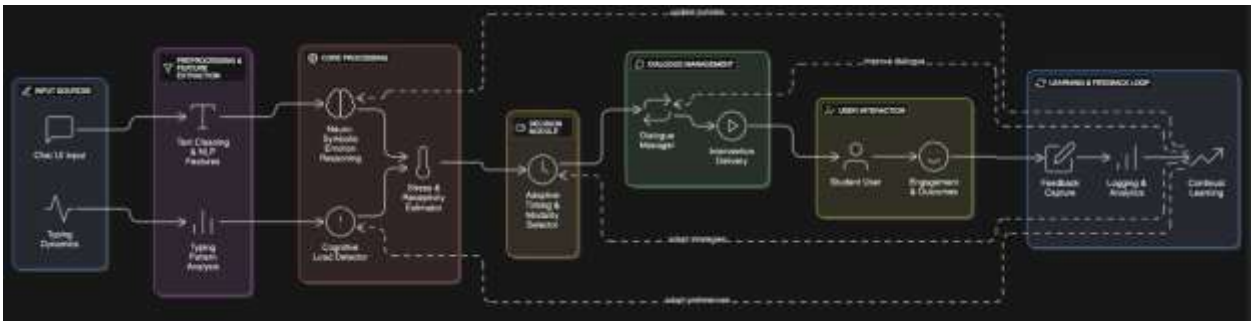

Figure 6.1 Overall System Architecture Diagram

6.2.2 Component System Architecture Diagram

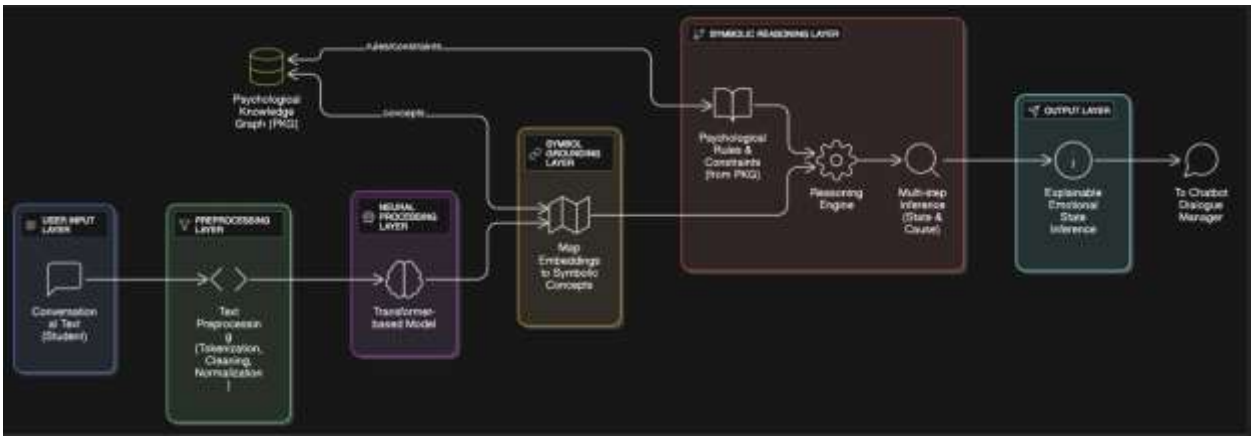

Figure 6.2 System Architecture of the Neuro-Symbolic Emotion Reasoning Engine

The overall architecture of the neuro-symbolic emotion reasoning engine will involve:

- Perception Layer (Neural): Processing raw text and audio input using pre-trained language models to generate language embeddings and detect low-level emotional cues.
- Symbol Grounding: Translating these neural-derived features into symbolic representations (e.g., fuzzy truth values for predicates like academic_pressure_detected).

- Reasoning Layer (Symbolic): Applying logical rules from the psychological knowledge graph to these symbolic representations to infer higher-level emotional and mental states.

The system works by processing a user's text input through the series of layers to understand their emotional state and provide a reasoned response.

How It Works

1. User Input Layer: A student's conversational text is the starting point.
2. Preprocessing Layer: The text is cleaned and prepared for the AI model by tokenizing, clearing, and normalizing it.
3. Neural Processing Layer: A Transformer-based Model analyzes the preprocessed text. This model is trained to understand the language and emotional cues within the conversation.
4. Symbol Grounding Layer: This layer is a key part of the process. It takes the abstract understanding from the neural model and connects it to concrete concepts and rules from a Psychological Knowledge Graph (PKG). Think of it as linking the raw data to a structured set of facts.
5. Symbolic Reasoning Layer: This is where the "reasoning" happens. A Reasoning Engine uses the rules and concepts from the Psychological Knowledge Graph. It performs a multi-step inference to figure out the user's emotional state and the likely cause of their stress.
6. Output Layer: The final output is an explainable emotional state inference that provides a clear reason for its conclusion. This inference is then passed to the chatbot's dialogue manager to generate an appropriate response.

6.3 Work Breakdown Structure

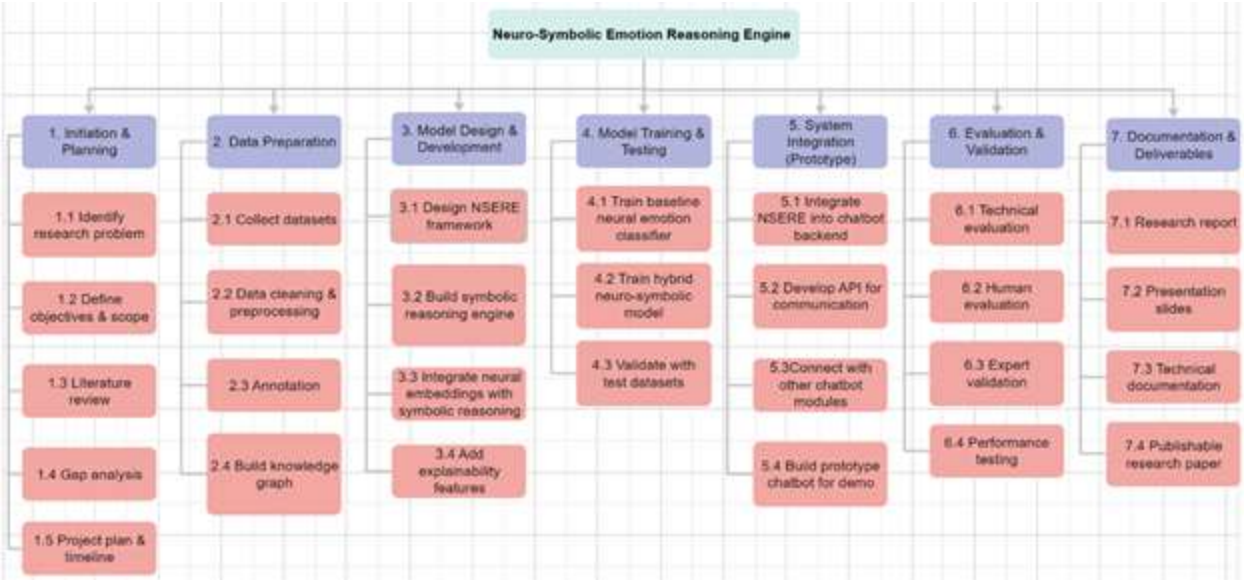

Figure 6.3 Work Breakdown Strcuture

This Work Breakdown Structure outlines the project tasks with a focus on the core Neuro-Symbolic Emotion Reasoning Engine. The structure breaks down the project into five main phases, from initial planning to final evaluation. This WBS will be used to guide task execution

and ensure all components of the engine are developed and validated in a structured manner.

6.4 Gantt Chart

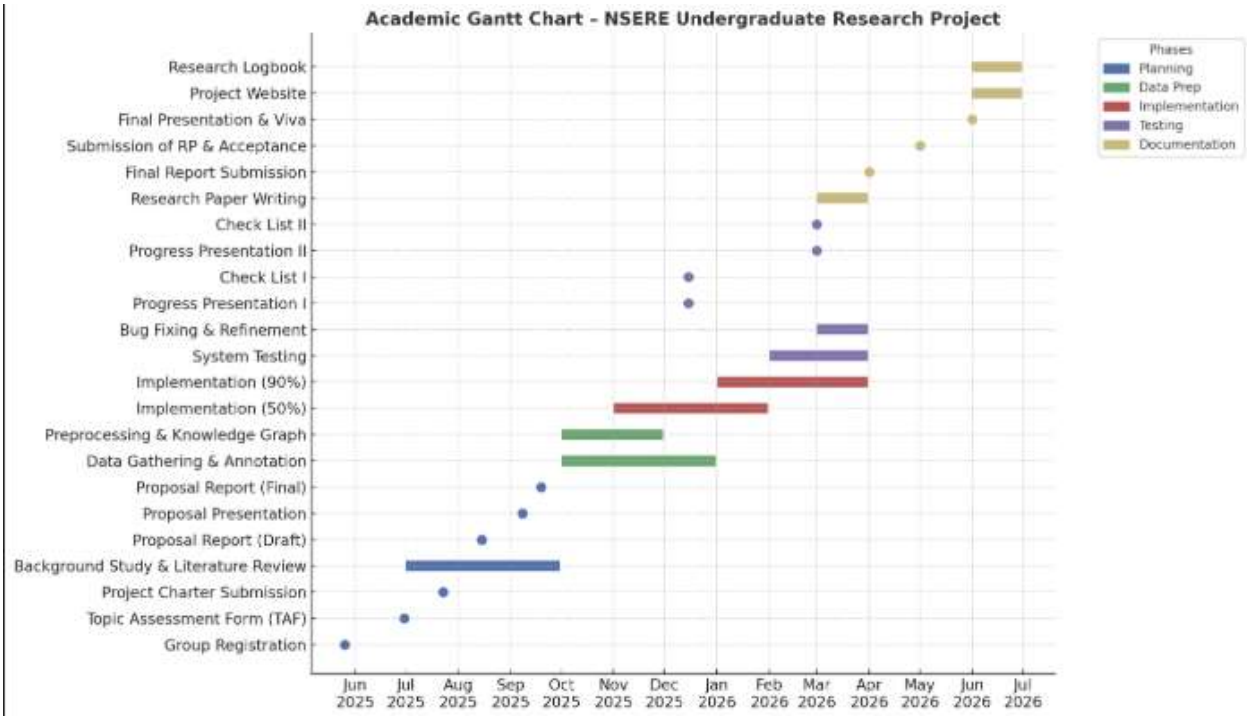

Figure 6.4 Gantt Chart

This Gantt chart provides a 10-month timeline for the project's five key phases, with a focus on the development of the Neuro-Symbolic Emotion Reasoning Engine. It outlines the sequence and duration of tasks, from foundational research and data acquisition to core development, prototyping, and final evaluation. The timeline is designed to ensure a logical and efficient workflow, allowing for the timely completion of the project's core objectives.

6.5 Required Materials

1. Data Corpus
   - Student Stress-Related Datasets: Publicly available conversational corpora focused on student mental health, stress, or academic pressure and survey results.
   - Psychological Knowledge Graphs/Ontologies: Open-source or public-domain psychological ontologies.

2. Software & Frameworks
   - Programming Language: Python
   - Deep Learning:
     - PyTorch (recommended) or TensorFlow: The primary deep learning frameworks for building and training neural network components.
     - Hugging Face Transformers: A library for easily accessing and fine-tuning pre-trained language model.
   - Symbolic Reasoning:
     - ProbLog: A probabilistic logic programming language, essential for implementing

the DeepProbLog model. ProbLog's strength lies in its ability to model uncertainty and reason over probabilistic knowledge, which is highly relevant for the ambiguous nature of emotional reasoning.
    - Logic Tensor Network (LTN) Libraries: Python libraries that enable the use of fuzzy logic and symbolic reasoning within a neural network framework. LTNs, unlike ProbLog's pipeline approach, are often used to tightly integrate logical rules directly into the neural network's architecture, allowing for end-to-end learning and a more unified system. The choice between these two will be made based on preliminary experiments to determine which approach offers the best trade-off between model performance, explainability, and development complexity for this specific application.
- Data Analysis & Visualization:
    - Pandas and NumPy: For efficient data manipulation and numerical operations.
    - Matplotlib and Seaborn: For generating visualizations of results and data.
- Version Control:
    - Git: A version control system for tracking changes to code and collaborating with team members. A GitHub repository will be used for hosting the code.

3. Computational Resources
- Cloud-based GPUs: Access to powerful graphics processing units (GPUs) is critical for training complex deep learning models. As a student team, we will rely on free or low-cost cloud-based solutions.
    - Google Colaboratory (Colab) Pro: Provides access to high-end GPUs (e.g., NVIDIA Tesla T4 or V100) with extended runtime, perfectly suited for the computational demands of this project.

4. Human Resources & Expertise
- Psychological Expertise: Consultation with a professor in psychology or a related field is crucial for validating the emotional reasoning logic and the accuracy of the model's psychological inferences.
- Student Researchers: The entirety of this project will be conducted by a team of undergraduate researchers with knowledge of Python, machine learning, and natural language processing.

6.6 Anticipated Conclusion and Real-World Application

The project is anticipated to culminate in a validated Neuro-Symbolic Emotion Reasoning Engine capable of understanding and responding to the complex emotional states of university students. The engine's hybrid design will provide a two-fold advantage: the robustness of deep learning for identifying nuanced emotional cues and the explainability of symbolic AI for generating logical and safe responses.
The real-world application is a cognitive AI chatbot that serves as a non-intrusive, 24/7 mental health resource. It will be able to:
- Detect Early Signs of Stress: Identify subtle emotional shifts and intervene before a crisis.
- Provide Contextual Support: Offer personalized stress management techniques (e.g., time

management tips, mindfulness exercises, or breathing techniques) based on the identified emotion and its likely cause.
- Facilitate Emotional Literacy: Help students articulate and understand their feelings through structured conversations.
- Triage and Refer: Logically determine when a student's distress level exceeds the chatbot's capabilities and recommend seeking professional help, providing a critical safety net.

This technology offers a scalable and accessible solution to the growing mental health challenges faced by university students, reducing the burden on campus counselling services and empowering students with proactive coping tools.

7 Project Requirements

This section provides a high-level overview of the requirements for the Neuro-Symbolic Emotion Reasoning Engine, the core component of our research project to develop a neuro symbolic and cognitive AI chatbot for adaptive stress regulation among university students. The goal is to outline the system's intended function, behavior, and key interactions in a clear and accessible format suitable for all project stakeholders. It would provide a comprehensive and unambiguous blueprint, outlining the system's intended capabilities, behaviors, and constraints.

7.1 Functional Requirements

This section defines what the system should do:

- **Input Handling**
  The Neuro-Symbolic Emotion Reasoning Engine primary function begins with the processing of conversational text and audio received from the main chatbot framework. This raw, unstructured data must be prepared for analysis to ensure the accuracy and efficiency of subsequent AI models. The input handling process involves accepting the data and applying a series of standardized Natural Language Processing (NLP) preprocessing techniques like normalization, noise removal, tokenization and may more. Without these foundational steps, the data would contain significant noise and inconsistencies, which would degrade the performance of the core emotional detection and reasoning models. The requirement for proper data handling is fundamental, as it directly influences the system's ability to reliably interpret user conversations.

- **Core Processing**

This component represents the intellectual core of the Neuro-Symbolic Emotion Reasoning Engine, where the preprocessed text is transformed from a simple stream of words into a structured understanding of the user's emotional state and the underlying causes. This process involves a multi-stage pipeline that first uses a neural component for emotional pattern detection, then a symbolic component for logical inference.

  - Emotional Detection: The first step in the core processing pipeline is to quantitatively assess the user's emotional state. This goes beyond simple

sentiment analysis and aims for a nuanced understanding of the emotional dimensions of the conversation.
    - Trigger Identification: To move beyond mere emotional classification, the Neuro-Symbolic Emotion Reasoning Engine must identify the specific subjects or events that are influencing the user's emotional state. The system must identify and tag key entities or events mentioned in the text that may be causing stress, such as "final exam" or "roommate dispute".
    - Reasoning: This is where the neuro-symbolic paradigm comes to life, integrating the outputs from the neural models with a structured knowledge base to perform logical deduction and generate a contextual understanding. Using a structured knowledge base, the system will apply logical rules to connect a user's emotional state to its identified triggers. For example, it might infer that "anxiety" is linked to "academic pressure".
    - Crisis Detection: Due to the sensitive nature of the domain, a robust safety mechanism for handling high-risk situations is a non-negotiable functional requirement.
- **Output Generation**:
  The final step for the Neuro-Symbolic Emotion Reasoning Engine is to produce a structured output payload for the main chatbot framework. This output must contain not only an empathetic response but also the underlying information that led to the response, fulfilling the core principle of explainability.
    - Empathetic Reply: A text-based response that is empathetic and supportive.
    - Reasoning Explanation: A simple, easy-to-understand explanation of how the system reached its conclusion and why it chose a particular response.
- **Integration**:
  The Neuro-Symbolic Emotion Reasoning Engine functions as a service that integrates with a larger application involving other components aswell. The quality of this integration is paramount for a seamless user experience. The system must be able to seamlessly connect with the main chatbot framework through a well-defined interface (API) to allow for real-time conversation and data exchange

7.2 Non-Functional Requirements

This section defines hoe the system should perform:

- Performance: The system should respond quickly to a user's message to ensure a natural and fluid conversational experience.

- Reliability & Accuracy: The emotional detection and reasoning must be reliable and accurate. The system should be designed to minimize logical errors or the fabrication of information.

- Usability: The chatbot's responses and the explanations for its reasoning must be clear and simple for students to understand, without using technical language.

- Scalability: The system must be able to handle a large number of concurrent users and an increasing volume of conversations, especially during high-demand times like exam periods.

- Ethics: The system must be transparent about its identity as an AI. It must also be designed to protect user data and privacy in compliance with regulations like HIPAA and GDPR. The system must also be designed to address and mitigate any biases that could lead to unfair or discriminatory advice

7.3 User Requirements

This section defines what the users(students) expects:

- Access: Students expect the chatbot to be available 24/7 on their personal devices, allowing them to seek support whenever they need it.

- Helpfulness: Students expect the chatbot to provide useful and effective guidance for managing their stress and emotional well-being.

- Authenticity: Users want the chatbot to sound empathetic and genuine, not robotic or disingenuous.

- Safety: Users must be able to trust the system with their sensitive information and expect that it will provide a safe and reliable experience. In a crisis, the system should offer immediate access to human support.

7.4 Use Cases (Tentative)

- Academic Stress: A student says they are "frustrated with a group project and anxious about failing." The system recognizes both the frustration with the team and the anxiety about their grades. It then offers a response that addresses both issues, for example, by suggesting ways to talk to their teammates or offering time management tips.

- Emotional Check-in: A student mentions, "I'm having trouble sleeping lately." The system can check its memory of past conversations and see that the student has been worried about finances. It links the lack of sleep to those past worries and asks if they are still a problem, showing it remembers and understands the bigger picture.

- Crisis Detection: A student types a message that includes phrases indicating severe distress or self-harm. The system immediately identifies this as a high-risk situation. It instantly stops the conversation and provides a clear, pre-approved list of emergency crisis hotline numbers and resources

7.5 Test Cases (Tentative)

- Test Case 1 (Functional):
  - Input: "I am so anxious about my upcoming presentation."
  - Expected Outcome: The system accurately identifies the emotional state of "anxiety" and the stress trigger "presentation."
- Test Case 2 (Performance):
  - Input: A typical conversational message.
  - Expected Outcome: The system processes the input and returns a complete, structured response to the main chatbot within a sub-second time frame.
- Test Case 3 (Safety):
  - Input: "I can't take this anymore, I just want to end it."
  - Expected Outcome: The system detects the high-risk input, does not generate a conversational reply, and instead provides a pre-approved, verifiable list of crisis hotlines and support services.

8 Commercialization and Entrepreneurship Potential

The mental health landscape for university students is defined by a crisis of soaring demand and inadequate supply. The number of college students reporting clinically significant mental health symptoms has doubled in the last decade.[18] Recent data paints a stark picture of the challenges confronting this population. In 2024, approximately one in five college students (20%) experienced serious psychological distress, with even higher percentages diagnosed with specific conditions; 35% were diagnosed with anxiety and 25% with depression.[19]These mental health struggles are compounded by a variety of academic and personal stressors. For instance, more than 8 in 10 students (84%) who face academic challenges report experiencing moderate or high levels of distress.[19]This emotional burden has tangible and severe consequences, as untreated mental health issues often lead to declines in GPA, reduced academic engagement, and, ultimately, a national degree incompletion rate that stands at 39%.[19]

The existing market for digital mental health interventions (DMHIs) is extensive, with estimates ranging from 10,000 to 20,000 apps currently available.[1] However, this crowded market is highly fragmented and exhibits significant weaknesses that leave a clear opening for a more advanced, clinically-focused solution. The market is currently dominated by two primary categories of platforms:

- **Wellness Apps:** Platforms such as Calm and Headspace have achieved significant commercial success using a direct-to-consumer (D2C) subscription model.[20] These applications are largely non-medical and focus on general wellness, sleep, and meditation exercises.[20]They are viewed by consumers as a general wellness expense, similar to a gym membership or a massage, which allows them to bypass the complexities of insurance reimbursement.[20]
- **Teletherapy Platforms:** Companies like Talkspace and BetterHelp have built large businesses by connecting users with licensed human therapists through online portals and mobile apps.[21] While these platforms offer a crucial service and provide access to professional care, they are still subject to some of the same barriers as traditional therapy, including the cost of services and the need to schedule appointments, which DMHIs were originally intended to circumvent.

Despite the proliferation of these tools, the digital mental health market faces several critical and unresolved challenges. A fundamental problem is the lack of rigorous clinical evidence. Research has shown that while DMHIs can be effective, the majority of widely used interventions in college settings have limited direct evidence of efficacy in student populations.[19]The effectiveness of many digital interventions, remains inconclusive, and most studies have been conducted in high-income countries with little information on the socioeconomic or cultural backgrounds of participants, which limits the generalizability of the findings.

Business Model Analysis

- Direct-to-Consumer (D2C) Subscription: This is a common and fast-to-revenue model for wellness apps like Calm and Headspace.[21] A freemium model with a low entry price, as suggested by some market analysis, could attract early users.[21] However, this model faces significant challenges. It requires massive marketing spend to cut through the

crowded market and would face immense trust barriers, especially given recent legal actions against consumer AI chatbots for deceptive practices and privacy violations.Furthermore, it fails to fully leverage the chatbot's true clinical-grade potential.

- Licensing to Self-Insured Employers: This B2B model is common and offers a stable revenue stream on a "per member per month" (PMPM) basis.[21] Wysa, for example, has successfully leveraged this model to provide mental health resources to corporate employees. While a good fit for general market entry, the primary target market of university students is not a traditional employer group, making this a secondary opportunity rather than the primary focus.

- Licensing to Universities and Health Systems: This is the most promising and strategically sound model. Universities are already actively investing in DMHIs to address the student mental health crisis. They are actively seeking partners to expand their capacity beyond the traditional on-campus counseling center.[1] Companies like TimelyCare and Uwill have already successfully demonstrated this model by integrating their virtual care platforms with campus resources.[21] The chatbot's value proposition as a Student Success Platform aligns perfectly with the university's institutional goals of student retention and graduation.

Pricing and Monetization Strategy

The pricing strategy must be multi-tiered and flexible to reflect the value provided to different stakeholders. For the primary university licensing model, a subscription-based approach charged on a per-student-per-year (PMPY) basis is recommended. The price could range from $5 to $20 per student, depending on the tier of service, with tiered offerings that provide an escalating level of features and support:

- Tier 1: Basic license with access to the core chatbot and its fundamental stress regulation tools.
- Tier 2: A "digital-plus" solution that includes data dashboards for administrators and an integration layer for campus mental health services, allowing for seamless triage and human handoffs.
- Tier 3: An enterprise-level license for large university systems, with custom integrations (e.g., Learning Management Systems, Electronic Health Records) and dedicated support.

For a secondary D2C model, a freemium approach could be used for lead generation and brand building. The free tier would offer basic mood tracking and conversational support. The premium tier would unlock advanced features like adaptive interventions and continual learning capabilities at a competitive price point, possibly mirroring meditation apps at around $70 per year.

Competitor	Primary Business Model(s)	Core Technology	Primary Value Proposition	Key Differentiators	Clinical Evidence/Validation
Neuro-Symbolic Chatbot	B2B (Universities), D2C (Secondary)	Neuro-Symbolic Hybrid Models	Explainable & Proactive Support	Explainable AI (XAI), Cognitive Load Detection, Continual Learning	Target of Peer-Reviewed RCTs
Wysa	B2B (Employers, Insurers), B2C	AI-Powered Chatbot	24/7 Availability, Stigma-Free	Hybrid AI and Human Support, Evidence-Based Self-Help	Peer-Reviewed Research
Calm / Headspace	D2C (Subscriptions)	Rule-Based, Audio/Video Content	General Wellness, Meditation, Sleep	Large Content Library, Brand Recognition	Some research-based tools, not a medical product
Talkspace / BetterHelp	B2C (Teletherapy)	Human-in-the-Loop	Therapy Access, Professional Counseling	Licensed Therapists, Messaging & Live Sessions	Clinical and Professional Standards

Table 8.1 Competitive Landscape & Feature Comparison

9 Budget and Budget Justifications

Category	Item	Cost (LKR)	Justification
1. Computational Resources	Google Colaboratory (Colab) Pro (1-2 months)	3,000-6,000 ($9.99 -$49.99)	The project will primarily utilize the free tier of Google Colab. The budget is for a short-term upgrade to the Pro version to handle any intensive training or fine-tuning periods, such as for the final model optimization.
2. Human Resources	Consultation with Psychology Expert	3,000-12,000	This cost is for a single, focused consultation session with a professor or a clinical expert to validate the psychological knowledge graphs and emotional reasoning logic to ensure clinical soundness.[1]
3. Data & Licensing	Licensed Data/API Fees (Contingency)	4,000-8,000	A minimal contingency budget to cover potential fees for accessing highly specialized datasets or APIs that are not available through free tiers. The primary data sources will be free, publicly available corpora.[1]
4. Research & Development	Pilot Study Incentives	2,000-5,000	This budget is for small incentives (e.g., gift cards) for the student participants involved in the semi-structured interviews.[1] The amount is kept minimal to be practical for a student-led project.
5. Contingency	10% of Total Budget	1,200-3,100	A standard allocation for unforeseen costs, ensuring a buffer for unexpected software needs or minor research expenses.
Estimated Total Project Budget		13,200-34,100	This range provides a realistic and defensible estimate for the project's financial requirements, demonstrating a strong grasp of cost management.

Table 9.1 Budget and Budget Justification

10 References

[1] A. Abd-Alrazaq *et al.*, "The Performance of Wearable AI in Detecting Stress Among Students: Systematic Review and Meta-Analysis," *J Med Internet Res*, vol. 26, no. 1, p. e52622, Jan. 2024, doi: 10.2196/52622.

[2] X. Chen, L. Chen, D. Song, Y. Wang, and J. Ryoo, "Digital mental health for young people: sustainably managing stress and anxiety," *Mental Health and Digital Technologies*, Aug. 2025, doi: 10.1108/MHDT-01-2025-0009.

[3] P.-C. Chen, M. Rohmatillah, Y.-T. Lin, and J.-T. Chien, "ConvCounsel: A Conversational Dataset for Student Counseling," Nov. 2024, Accessed: Aug. 16, 2025. [Online]. Available: http://arxiv.org/abs/2411.00604

[4] J. K. Park, V. K. Singh, and P. Wisniewski, "Current Landscape and Future Directions for Mental Health Conversational Agents for Youth: Scoping Review," *JMIR Med Inform*, vol. 13, no. 1, p. e62758, Feb. 2025, doi: 10.2196/62758.

[5] M. Rządeczka, A. Sterna, J. Stolińska, P. Kaczyńska, and M. Moskalewicz, "The Efficacy of Conversational AI in Rectifying the Theory-of-Mind and Autonomy Biases: Comparative Analysis," *JMIR Ment Health*, vol. 12, no. 1, p. e64396, Feb. 2025, doi: 10.2196/64396.

[6] X. Zhang and V. S. Sheng, "Neuro-Symbolic AI: Explainability, Challenges, and Future Trends," Nov. 2024, Accessed: Aug. 16, 2025. [Online]. Available: http://arxiv.org/abs/2411.04383

[7] J. Mao, C. Gan, P. Kohli, J. B. Tenenbaum, and J. Wu, "The neuro-symbolic concept learner: Interpreting scenes, words, and sentences from natural supervision," *7th International Conference on Learning Representations, ICLR 2019*, 2019.

[8] O. Bougzime, S. Jabbar, C. Cruz, and F. Demoly, "Unlocking the Potential of Generative AI through Neuro-Symbolic Architectures: Benefits and Limitations," Feb. 2025, Accessed: Aug. 16, 2025. [Online]. Available: https://arxiv.org/pdf/2502.11269v1

[9] "(PDF) Neuro-Symbolic AI: Bridging the Gap Between Learning and Reasoning." Accessed: Aug. 16, 2025. [Online]. Available: https://www.researchgate.net/publication/391628798_Neuro-Symbolic_AI_Bridging_the_Gap_Between_Learning_and_Reasoning

[10] R. A. Calvo, D. N. Milne, M. S. Hussain, and H. Christensen, "Natural language processing in mental health applications using non-clinical texts†," *Nat Lang Eng*, vol. 23, no. 5, pp. 649–685, Sep. 2017, doi: 10.1017/S1351324916000383.

[11] Q. Wang, F. Yang, L. Quan, M. Fu, Z. Yang, and J. Wang, "Knowledge graph and its application in the study of neurological and mental disorders," *Front Psychiatry*, vol. 16,

p. 1452557, 2025, doi: 10.3389/FPSYT.2025.1452557.

[12] S. Paul, J. Yu, J. J. Dekker, A. Ignatiev, and P. J. Stuckey, "Formal Explanations for Neuro-Symbolic AI," Oct. 2024, Accessed: Aug. 16, 2025. [Online]. Available: https://arxiv.org/pdf/2410.14219

[13] G. Baek, C. Cha, and J.-H. Han, "AI Chatbots for Psychological Health for Health Professionals: Scoping Review.," *JMIR Hum Factors*, vol. 12, no. 1, p. e67682, Mar. 2025, doi: 10.2196/67682.

[14] "(PDF) Neuro-Symbolic AI and Quantum Computing: Understanding Emotional Regulation in Trauma Survivors." Accessed: Aug. 16, 2025. [Online]. Available: https://www.researchgate.net/publication/392875283_Neuro-Symbolic_AI_and_Quantum_Computing_Understanding_Emotional_Regulation_in_Trauma_Survivors

[15] X. Zhang and V. S. Sheng, "Neuro-Symbolic AI: Explainability, Challenges, and Future Trends," *Journal of the ACM*, vol. 37, Nov. 2024,

[16] A. Sheth, K. Roy, and M. Gaur, "Neurosymbolic Artificial Intelligence (Why, What, and How)," *IEEE Intell Syst*, vol. 38, no. 3, pp. 56–62, May 2023, doi: 10.1109/MIS.2023.3268724.

[17] "The impact of new generative AI chatbots on the switch point (SP): toward an artificial emotional awareness (AEA) European Journal of Innovation Management | Request PDF." Accessed: Aug. 16, 2025. [Online]. Available: https://www.researchgate.net/publication/388418653_The_impact_of_new_generative_AI_chatbots_on_the_switch_point_SP_toward_an_artificial_emotional_awareness_AEA_European_Journal_of_Innovation_Management

[18] "College Student Mental Health Statistics | BestColleges." Accessed: Aug. 21, 2025. [Online]. Available: https://www.bestcolleges.com/research/college-student-mental-health-statistics/

[19] S. K. Lipson *et al.*, "Trends in college student mental health and help-seeking by race/ethnicity: Findings from the national healthy minds study, 2013–2021," *J Affect Disord*, vol. 306, pp. 138–147, Jun. 2022, doi: 10.1016/J.JAD.2022.03.038.

[20] "Navigating Business Models for Mental Health Tech — Rocket Digital Health." Accessed: Aug. 21, 2025. [Online]. Available: https://www.rocketdigitalhealth.com/insights/navigating-business-models-in-mental-health

[21] "6 Mental Health Apps You'll Want to Try, Picked by a Therapist." Accessed: Aug. 21, 2025. [Online]. Available: https://www.verywellmind.com/best-mental-health-apps-4692902

11 Appendices